

fn make_stable_truncate

Michael Shoemate

February 23, 2026

This proof resides in “**contrib**” because it has not completed the vetting process.

Proves soundness of `make_stable_truncate` in `mod.rs` at commit `f5bb719` (outdated¹).

1 Hoare Triple

Precondition

Caller Verified

None

Function

```
1 def make_stable_truncate(
2     input_domain: DslPlanDomain,
3     input_metric: FrameDistance[SymmetricIdDistance],
4     plan: DslPlan,
5 ) -> Transformation[
6     DslPlanDomain,
7     DslPlanDomain,
8     FrameDistance[SymmetricIdDistance],
9     FrameDistance[SymmetricIdDistance],
10 ]:
11     # the identifier is protected from changes, so we can use the identifier from the input
12     # metric
13     # instead of the identifier from the middle_metric to match truncations
14     input, truncations, truncation_bounds = match_truncations(
15         plan, input_metric[0].identifier
16     )
17
18     if truncations.is_empty():
19         return ValueError("failed to match truncation")
20
21     t_prior = input.make_stable(input_domain, input_metric)
22     middle_domain, middle_metric = t_prior.output_space()
23
24     for bound in truncation_bounds:
25         for key in bound.by:
26             # raises if the key is not infallible row-by-row
27             make_stable_expr(  #
28                 WildExprDomain(
29                     columns=middle_domain.series_domains,
30                     context=Context.RowByRow,
```

¹See new changes with git diff f5bb719..eb0a360 rust/src/transformations/make_stable_lazyframe/truncate/mod.rs

```

30         ),
31         PartitionDistance(middle_metric[0]),
32         key,
33     )
34
35     output_domain = middle_domain.clone()
36     for truncation in truncations: #
37         output_domain = truncate_domain(output_domain, truncation)
38
39     def function(plan: DslPlan) -> DslPlan:
40         for truncation in truncations:
41             match truncation:
42                 case Truncation.Filter(predicate):
43                     plan = DslPlan.Filter(
44                         input=plan,
45                         predicate=predicate,
46                     )
47
48                 case Truncation.GroupBy(keys, aggs):
49                     plan = DslPlan.GroupBy(
50                         input=plan,
51                         keys=keys,
52                         aggs=aggs,
53                         apply=None,
54                         maintain_order=False,
55                         options=GroupbyOptions.default(),
56                     )
57         return plan
58
59     def stability_map(id_bounds: Bounds) -> Bounds:
60         #
61         total_num_ids = id_bounds.get_bound({}).per_group
62
63         # each truncation is used to derive row bounds
64         new_bounds = []
65         for truncation_bound in truncation_bounds: #
66             # each truncation is used to derive row bounds
67             new_bounds.append(
68                 truncate_id_bound( #
69                     id_bounds.get_bound(truncation_bound.by), #
70                     truncation_bound,
71                     total_num_ids,
72                 )
73             )
74         return Bounds(new_bounds)
75
76     t_truncate = Transformation.new(
77         middle_domain,
78         output_domain,
79         Function.new(function),
80         middle_metric,
81         FrameDistance(SymmetricDistance),
82         StabilityMap.new_fallible(stability_map),
83     )
84     return t_prior >> t_truncate

```

Postcondition

Theorem 1.1. For every setting of the input parameters (`input_domain`, `input_metric`, `plan`) to `make_stable_truncate` such that the given preconditions hold,
`make_stable_truncate` raises an error (at compile time or run time) or returns a valid transformation. A valid transformation has the following properties:

1. (Data-independent runtime errors). For every pair of members x and x' in `input_domain`, `invoke(x)`

and `invoke`(x') either both return the same error or neither return an error.

2. (Appropriate output domain). For every member x in `input_domain`, `function`(x) is in `output_domain` or raises a data-independent runtime error.
3. (Stability guarantee). For every pair of members x and x' in `input_domain` and for every pair (d_{in}, d_{out}) , where d_{in} has the associated type for `input_metric` and d_{out} has the associated type for `output_metric`, if x, x' are d_{in} -close under `input_metric`, `stability_map`(d_{in}) does not raise an error, and `stability_map`(d_{in}) = d_{out} , then `function`(x), `function`(x') are d_{out} -close under `output_metric`.

Appropriate Output Domain. By line 26, the grouping keys are stable row-by-row transformations of the data, therefore the preconditions of 36 are satisfied. By the postcondition of 36, for every element x in `input_domain`, `function`(x) is in `output_domain` or raises a data-independent runtime exception. \square

Stability guarantee. By line 60, `total_num_ids` is the total number of ids an individual may contribute to a dataset:

$$d_{SymId}(\text{function}(x), \text{function}(x')) \leq \text{total_num_ids}. \quad (1)$$

By the postcondition of `match_truncations`, for each `truncation_bound` on line 65,

$$\begin{aligned} \max_{id} \|d_{Sym}(\text{function}(x)_{id,g}, \text{function}(x')_{id,g})\|_\infty &\leq \text{truncation.per_group}, \\ \max_{id} \|d_{Sym}(\text{function}(x)_{id,g}, \text{function}(x')_{id,g})\|_0 &\leq \text{truncation.num_groups}, \end{aligned}$$

where g denotes the group when partitioned by `truncation_bound.by`.

The preconditions of 68 are satisfied on line 69 (`id_bound.by` is equal to `truncation.by`), so by the postcondition of `truncate_id_bound`,

$$\begin{aligned} \|d_{Sym}(\text{function}(x)_g, \text{function}(x')_g)\|_\infty &\leq \text{row_bound.per_group}, \\ \|d_{Sym}(\text{function}(x)_g, \text{function}(x')_g)\|_0 &\leq \text{row_bound.num_groups}, \end{aligned}$$

where `row_bound` denotes the return value.

For each truncation on line 65, `truncate_id_bound` on line 68 computes upper bounds on the resulting distance between adjacent datasets. All acquired bounds are valid upper bounds on the distance between the two datasets, by the postcondition of `match_truncations`, that each truncation does not invalidate the truncation bounds of the previous truncations.

It is shown that for every pair of elements x, x' in `input_domain` and for every pair (d_{in}, d_{out}) , where d_{in} has the associated type for `input_metric` and d_{out} has the associated type for `output_metric`, if x, x' are d_{in} -close under `input_metric`, `stability_map`(d_{in}) does not raise an exception, and `stability_map`(d_{in}) $\leq d_{out}$, then `function`(x), `function`(x') are d_{out} -close under `output_metric`. \square